

# INFSCI 2935: Human-Robot Interaction

## Reinforcement Learning Project

### Vishruth Pradeep Reddy

[vpr8@pitt.edu](mailto:vpr8@pitt.edu)

Before we dive into the project, I would like to provide a brief description as to what each hyperparameter is and how they affect the model with change in values.

#### Hyper Parameter Description

##### 1. Buffer Size:

- Description: The size of the memory for past experiences.
- Purpose: Stores diverse experiences for learning.
- Increase: Captures more diverse experiences but increases memory usage.
- Decrease: Limits experience variety, potentially affecting learning quality.

##### 2. Learning Starts:

- Description: Delay before the model begins learning.
- Purpose: Gathers initial data for informed learning.
- Increase: Better data collection but delays learning.
- Decrease: Starts learning sooner but might begin with noisy or sparse data.

##### 3. Batch Size:

- Description: Number of experiences used in each learning update.
- Purpose: Affects learning stability and speed.
- Increase: More stable learning but slower training due to larger computations.
- Decrease: Faster training but might result in less stable learning due to noisier updates.

##### 4. Gamma:

- Description: Balances immediate vs. future rewards.
- Purpose: Determines the importance of future rewards.
- Increase: Prioritizes future rewards but might overlook immediate gains.
- Decrease: Focuses more on immediate rewards, potentially disregarding long-term gains.

##### 5. Train Frequency:

- Description: Frequency of updating the model's parameters.
- Purpose: Controls learning speed and parameter updates.
- Increase: Faster learning but might introduce more instability due to frequent updates.
- Decrease: Slower learning but could lead to more stable updates.

##### 6. Gradient Steps:

- Description: Number of optimization steps after each batch.
- Purpose: Affects convergence speed and stability.
- Increase: Faster convergence but might introduce instability due to larger updates.
- Decrease: Slower convergence but potentially more stable updates.

**7. Target Update Interval:**

- Description: Frequency of updating a secondary network.
- Purpose: Stabilizes learning by controlling update frequency.
- Increase: More stable learning but slower adaptation to changes.
- Decrease: Faster adaptation to changes but might introduce instability.

**8. Exploration Fraction:**

- Description: Determines exploration vs. exploitation.
- Purpose: Balances new strategy discovery and exploiting known ones.
- Increase: More exploration but might delay exploiting known strategies.
- Decrease: Exploits known strategies faster but might miss better ones.

**9. Verbose:**

- Description: Controls the level of output during training.
- Purpose: Provides information about the training process.
- Increase: More detailed information displayed during training.
- Decrease: Less information or no output during training.

**10. Training Steps:**

- Description: The total number of steps for training the model.
- Purpose: Sets the duration or extent of the training process.
- Increase: More training steps might lead to further learning and refinement.
- Decrease: Fewer steps might result in quicker but potentially less optimal learning.

## Algorithm 1: Initial DQN Model (as given in the boiler-plate code)

Survived 20+ seconds? NO

Hyperparameter	Description	Value
Network Architecture	Number of hidden layers	2
Neuron Count	Number of neurons in each hidden layer	[256, 256]
Learning rate	Rate of parameter updates	5e-4
Buffer size	Size of experience replay buffer	15000
Learning starts	Number of steps before learning begins	200
Batch size	Number of experiences per training iteration	32
Gamma	Discount factor for future rewards	0.8
Train frequency	Frequency of model training	1
Gradient steps	Number of optimization steps after each batch	1
Target update interval	Frequency of updating target network	50
Exploration fraction	Fraction of exploration during training	0.7
Verbose	Verbosity level during training	1
Training steps	Total steps for training	20,000

Model:

```
Initial DQN code ☆
File Edit View Insert Runtime Tools Help All
+ Code + Text
model = DQN('MlpPolicy', 'highway-fast-v0',
            policy_kwargs=dict(net_arch=[256, 256]),
            learning_rate=5e-4,
            buffer_size=15000,
            learning_starts=200,
            batch_size=32,
            gamma=0.8,
            train_freq=1,
            gradient_steps=1,
            target_update_interval=50,
            exploration_fraction=0.7,
            verbose=1,
            tensorboard_log='highway_dqn/')
model.learn(int(2e4))
```



## Final Training Stats:

```
Initial DQN code ☆
File Edit View Insert Runtime Tools Help All chana Comment
+ Code + Text
| n_updates | 19642 |
|-----|
| rollout/ |
| ep_len_mean | 20.8 |
| ep_rew_mean | 17.5 |
| exploration_rate | 0.05 |
| time/ |
| episodes | 1224 |
| fps | 16 |
| time_elapsed | 1218 |
| total_timesteps | 19920 |
| train/ |
| learning_rate | 0.0005 |
| loss | 0.104 |
| n_updates | 19719 |
|-----|
| rollout/ |
| ep_len_mean | 20.6 |
| ep_rew_mean | 17.2 |
| exploration_rate | 0.05 |
| time/ |
| episodes | 1228 |
| fps | 16 |
| time_elapsed | 1220 |
| total_timesteps | 19962 |
| train/ |
| learning_rate | 0.0005 |
| loss | 0.0234 |
| n_updates | 19761 |
|-----|
<stable_baselines3.dqn.dqn.DQN at 0x7ae5d4109e40>
```

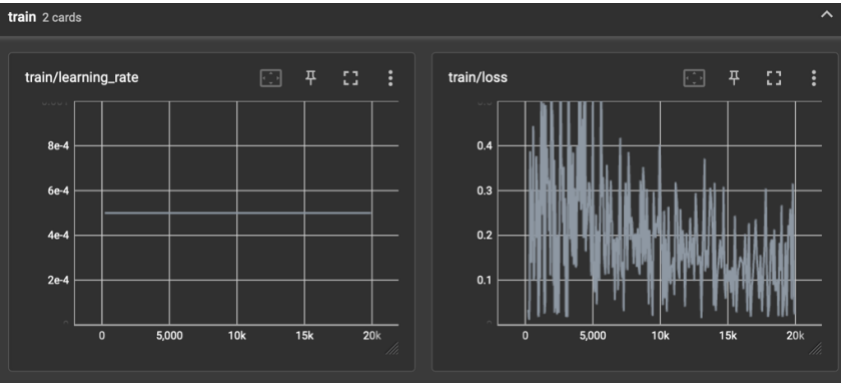
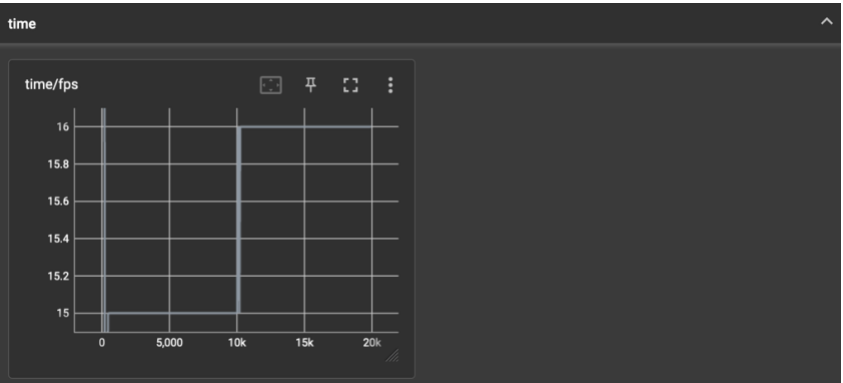
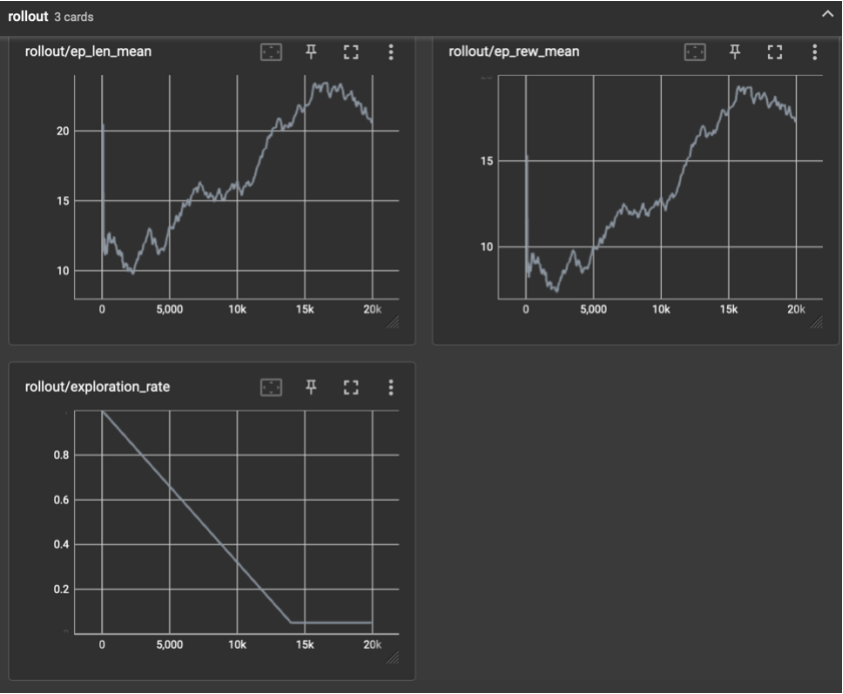
## Images:

The model could run without crashing for a maximum period of 13 seconds.

Image of the crash in 2 episodes, ran successfully in 1 episode:



Graphs:



## Algorithm 2: DQN Model (with low hyperparameter values for Learning Rate, Batch Size and Target Update Interval)

Survived 20+ seconds? YES

Hyperparameter	Description	Value
Network Architecture	Number of hidden layers	2
Neuron Count	Number of neurons in each hidden layer	[384, 384]
Learning rate	Rate of parameter updates	1e-4
Buffer size	Size of experience replay buffer	30000
Learning starts	Number of steps before learning starts	1000
Batch size	Number of experiences per training iteration	32
Gamma	Discount factor for future rewards	0.9
Train frequency	Frequency of model training	4
Gradient steps	Number of optimization steps after each batch	4
Target update interval	Frequency of updating target network	200
Exploration fraction	Fraction of exploration during training	0.5
Verbose	Level of output during training	1
Training steps	Total steps for training	30000

Model:

```
model_low = DQN(['MlpPolicy', 'highway-fast-v0',  
    policy_kwargs=dict(net_arch=[384, 384]), # Using 384 neurons in each hidden layer  
    learning_rate=1e-4, # Lowered learning rate to 1e-4  
    buffer_size=30000,  
    learning_starts=1000,  
    batch_size=32, # Reduced batch size to 32  
    gamma=0.9,  
    train_freq=4,  
    gradient_steps=4,  
    target_update_interval=200, # Increased target update interval to 200  
    exploration_fraction=0.5,  
    verbose=1,  
    tensorboard_log='highway_dqn_low/'])  
model_low.learn(total_timesteps=int(3e4))
```

### Final Training Stats:

Updated DQN low ☆

File Edit View Insert Runtime Tools Help All changes saved

+ Code + Text

learning_rate	0.0001
loss	0.275
n_updates	28692

---


rollout/	
ep_len_mean	26.9
ep_rew_mean	20.9
exploration_rate	0.05
time/	
episodes	1696
fps	12
time_elapsed	2334
total_timesteps	29788
train/	
learning_rate	0.0001
loss	0.199
n_updates	28784

---

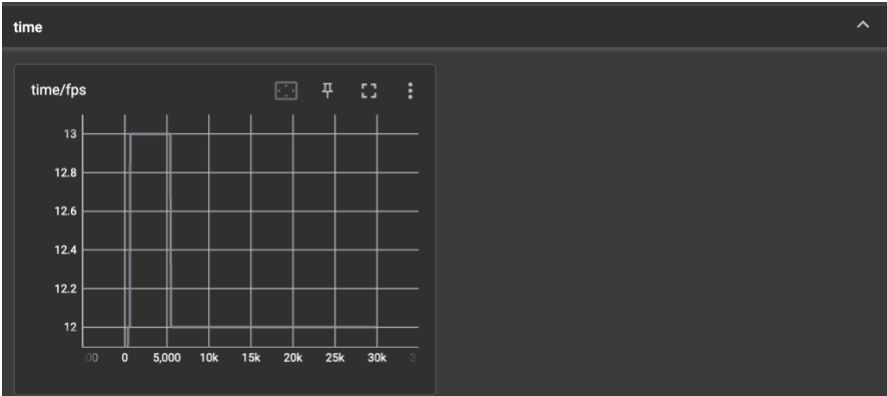
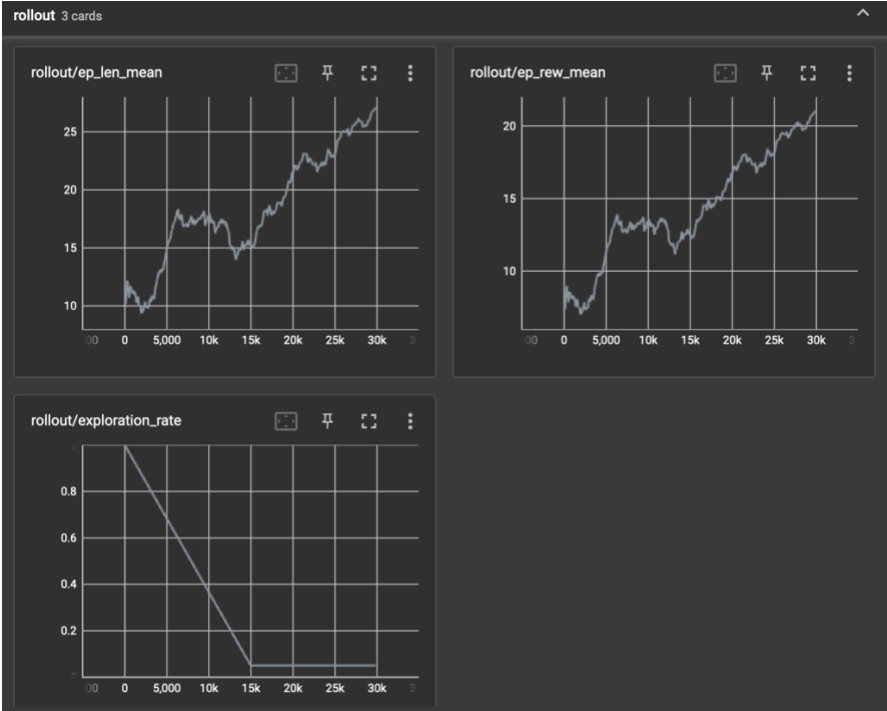
rollout/	
ep_len_mean	27.1
ep_rew_mean	21
exploration_rate	0.05
time/	
episodes	1700
fps	12
time_elapsed	2343
total_timesteps	29907
train/	
learning_rate	0.0001
loss	0.211
n_updates	28904

<stable\_baselines3.dqn.dqn.DQN at 0x7a36a3e0ac50>

Image (No signs for crashing in any episode, ran successfully in all 3 episodes):  
Travel Speed for car was relatively slow.



Graphs:





### Algorithm 3: (DQN with high hyperparameter values for Learning Rate, Batch Size and Target Update Interval)

Survived 20+ seconds? NO

Hyperparameter	Description	Value
Network Architecture	Number of hidden layers	2
Neuron Count	Number of neurons in each hidden layer	[384, 384]
Learning rate	Rate of parameter updates	5e-3
Buffer size	Size of experience replay buffer	30000
Learning starts	Number of steps before learning starts	1000
Batch size	Number of experiences per training iteration	128
Gamma	Discount factor for future rewards	0.9
Train frequency	Frequency of model training	4
Gradient steps	Number of optimization steps after each batch	4
Target update interval	Frequency of updating target network	500
Exploration fraction	Fraction of exploration during training	0.5
Verbose	Level of output during training	1
Training steps	Total steps for training	30000

Model:

```
model_high = DQN(['MlpPolicy', 'highway-fast-v0',
                  policy_kwargs=dict(net_arch=[384, 384]),
                  learning_rate=5e-3, # Increased learning rate to 5e-3
                  buffer_size=30000,
                  learning_starts=1000,
                  batch_size=128, # Increased batch size to 128
                  gamma=0.9,
                  train_freq=4,
                  gradient_steps=4,
                  target_update_interval=500, # Increased target update interval to 500
                  exploration_fraction=0.5,
                  verbose=1,
                  tensorboard_log='highway_dqn_high/'])
model_high.learn(total_timesteps=int(3e4))
```

## Final Training Stats:

Updated DQN high ☆

File Edit View Insert Runtime Tools Help Saving...

+ Code + Text



time/		
episodes	1840	
fps	11	
time_elapsed	2667	
total_timesteps	29771	
train/		
learning_rate	0.005	
loss	0.213	
n_updates	28768	

rollout/		
ep_len_mean	23.1	
ep_rew_mean	20.8	
exploration_rate	0.05	
time/		
episodes	1844	
fps	11	
time_elapsed	2676	
total_timesteps	29883	
train/		
learning_rate	0.005	
loss	0.146	
n_updates	28880	

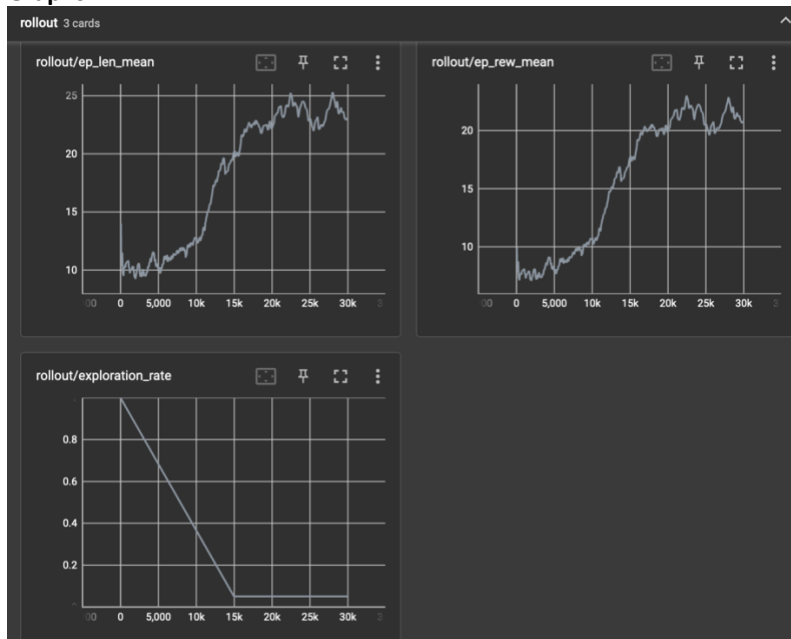
rollout/		
ep_len_mean	22.9	
ep_rew_mean	20.6	
exploration_rate	0.05	
time/		
episodes	1848	
fps	11	
time_elapsed	2685	
total_timesteps	29987	
train/		
learning_rate	0.005	
loss	0.191	
n_updates	28984	

<stable\_baselines3.dqn.dqn.DQN at 0x7eea175b7e20>

Image (Car crashed in 1 episode and ran successfully in 2 episodes):  
Travel Speed of cars was relatively fast.



### Graphs:



## Algorithm 4: (DQN with my own hyperparameter values for optimal run period of 20+ seconds)

Survived 20+ seconds? YES

Hyperparameter	Description	Value
Network Architecture	Number of hidden layers	2
Neuron Count	Number of neurons in each hidden layer	[256, 256]
Learning rate	Rate of parameter updates	1e-3
Buffer size	Size of experience replay buffer	20000
Learning starts	Number of steps before learning starts	1000
Batch size	Number of experiences per training iteration	32
Gamma	Discount factor for future rewards	0.95
Train frequency	Frequency of model training	4
Gradient steps	Number of optimization steps after each batch	4
Target update interval	Frequency of updating target network	100
Exploration fraction	Fraction of exploration during training	0.5
Verbose	Level of output during training	1
Training steps	Total steps for training	20000

Model:

```
Updated DQN ☆
File Edit View Insert Runtime Tools Help Saving...

Code + Text
The tensorboard extension is already loaded. To reload it, use:
[ ] %reload_ext tensorboard
Requirement already satisfied: tensorboardx in /usr/local/lib/python3.10
Requirement already satisfied: gym in /usr/local/lib/python3.10/dist-pac
Requirement already satisfied: pyvirtualdisplay in /usr/local/lib/python

model = DQN('MlpPolicy', 'highway-fast-v0',
            policy_kwargs=dict(net_arch=[256, 256]),
            learning_rate=1e-3, # Higher learning rate
            buffer_size=20000,
            learning_starts=1000,
            batch_size=32,
            gamma=0.95, # Slightly higher gamma
            exploration_fraction=0.5, # Reduced exploration
            train_freq=4, # Train less frequently
            gradient_steps=4, # Increase gradient steps
            target_update_interval=100, # Less frequent target updates
            verbose=1,
            tensorboard_log='highway_dqn/')

model.learn(int(2e4))
```

## Final Stats of Training:

Updated DQN ☆

File Edit View Insert Runtime Tools Help Last saved at 1:28 PM

Code + Text

▶

🔗

time/	episodes	916	
	fps	13	
	time_elapsed	1515	
	total_timesteps	19743	
train/	learning_rate	0.001	
	loss	0.343	
	n_updates	18740	

-----

rollout/	ep_len_mean	27.9	
	ep_rew_mean	21.3	
	exploration_rate	0.05	
time/	episodes	920	
	fps	13	
	time_elapsed	1522	
	total_timesteps	19852	
train/	learning_rate	0.001	
	loss	0.262	
	n_updates	18848	

-----

rollout/	ep_len_mean	27.9	
	ep_rew_mean	21.3	
	exploration_rate	0.05	
time/	episodes	924	
	fps	13	
	time_elapsed	1532	
	total_timesteps	19969	
train/	learning_rate	0.001	
	loss	0.458	
	n_updates	18968	

-----

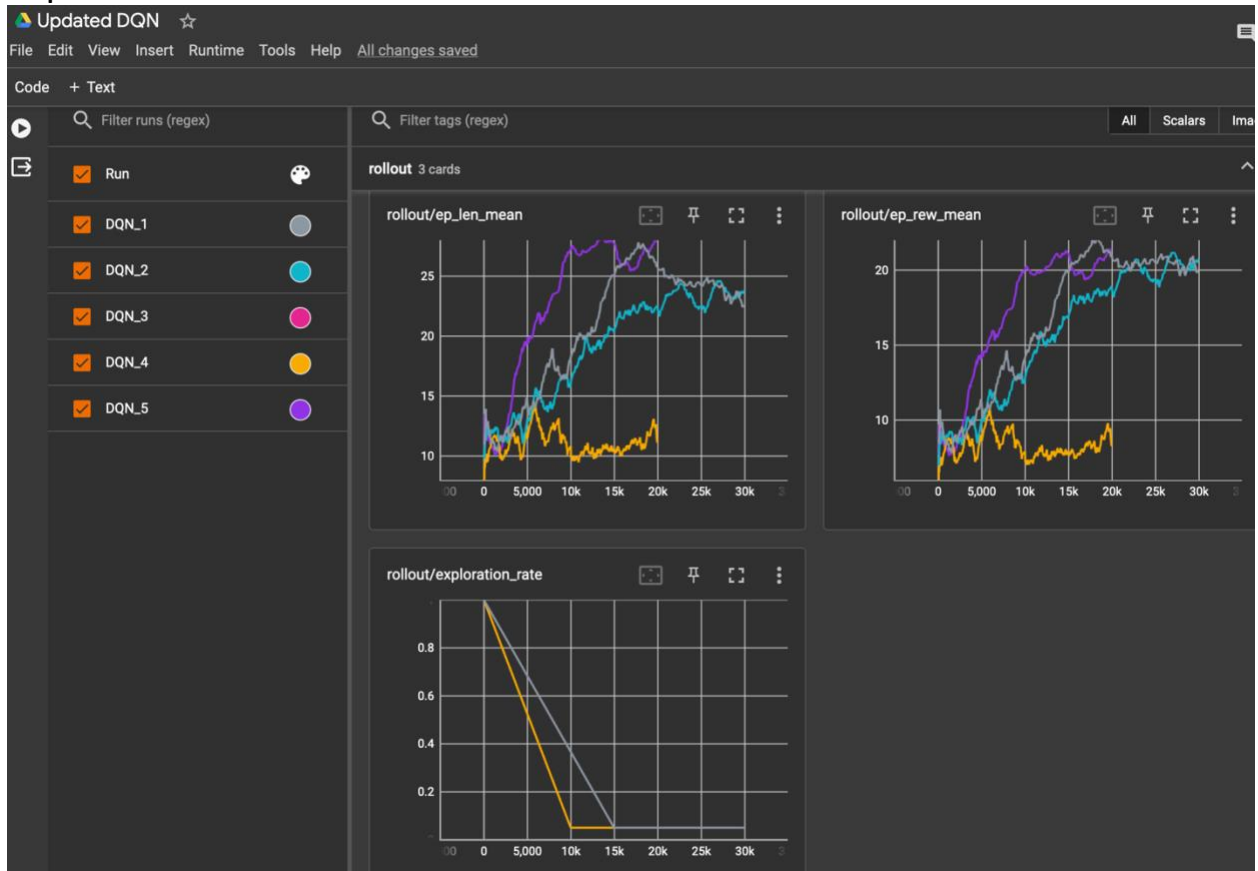
<stable\_baselines3.dqn.dqn.DQN at 0x7c8a3f0c0dc0>

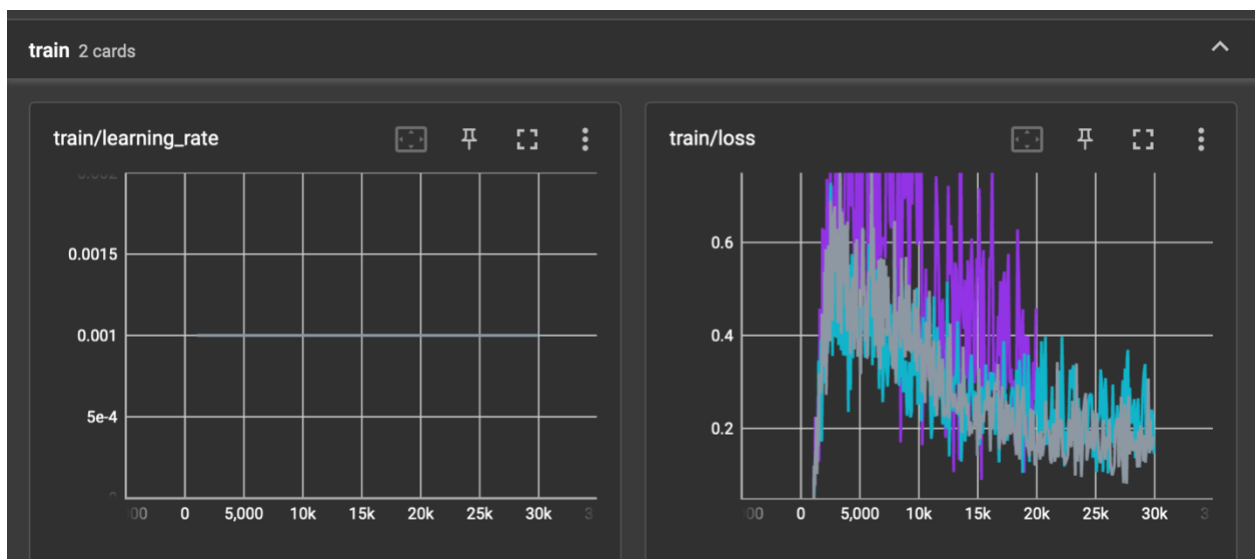
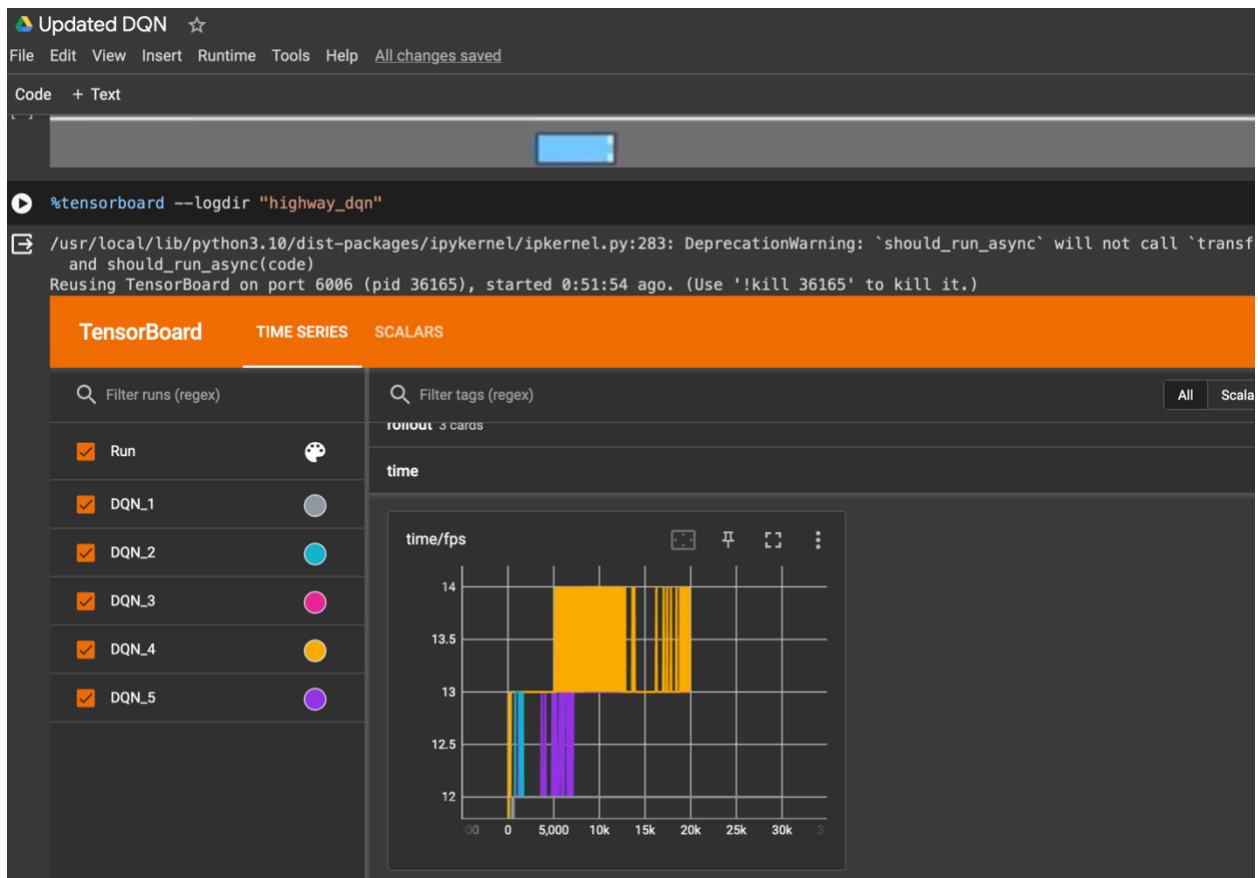
Image:

No signs of crashing in all episodes.



Graphs:





# Algorithm 5: PPO (Proximal Policy Optimization Algorithm)

Survived 20+ seconds? YES

Hyperparameter	Description	Value
Network Architecture	Number of hidden layers	4
Neuron Count	Number of neurons in each hidden layer	[256, 256, 256, 256]
Batch size	Number of steps to optimize at each iteration	128
n_epochs	Number of optimization epochs	5
Learning rate	Rate of learning for the optimizer	1e-4
Gamma	Discount factor for future rewards	0.9
Verbose	Level of output during training	2
Training steps	Total steps for training	20000

Model:

```
model = PPO(["MlpPolicy", "highway-fast-v0",
            policy_kwargs=dict(net_arch=[dict(pi=[256, 256], vf=[256, 256])]),
            n_steps=256,
            batch_size=128,
            n_epochs=5, |
            learning_rate=1e-4,
            gamma=0.9,
            verbose=2,
            tensorboard_log="highway_ppo/", )
model.learn(int(2e4))
```

iterations	77	
time_elapsed	1115	
total_timesteps	19712	
train/		
approx_kl	0.0015526903	



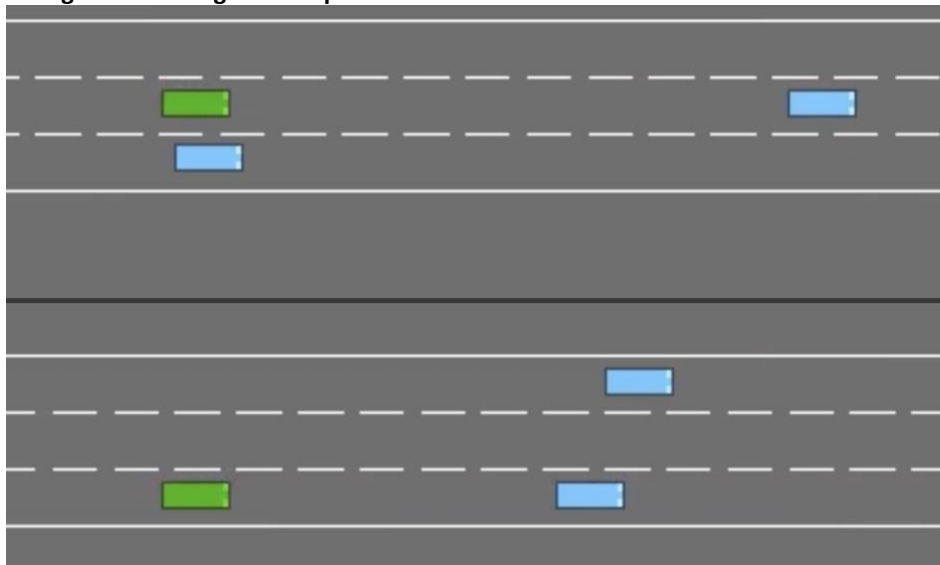
## Final Training Stats:

```
PPO.ipynb ☆
File Edit View Insert Runtime Tools Help All changes saved
+ Code + Text

▶ | ep_reward_mean | 19.4 |
  | time/          |     |
  | fps            | 17   |
  | iterations     | 78   |
  | time_elapsed   | 1129 |
  | total_timesteps | 19968 |
  | train/         |     |
  | approx_kl      | 0.00043736212 |
  | clip_fraction  | 0     |
  | clip_range     | 0.2   |
  | entropy_loss   | -0.82 |
  | explained_variance | 0.118 |
  | learning_rate  | 0.0001 |
  | loss           | 0.328 |
  | n_updates      | 385   |
  | policy_gradient_loss | -0.00086 |
  | value_loss     | 0.583 |
  |-----|-----|
  | rollout/       |     |
  | ep_len_mean    | 28.2  |
  | ep_reward_mean | 19.9  |
  | time/          |     |
  | fps            | 17    |
  | iterations     | 79    |
  | time_elapsed   | 1144  |
  | total_timesteps | 20224 |
  | train/         |     |
  | approx_kl      | 0.0001614329 |
  | clip_fraction  | 0     |
  | clip_range     | 0.2   |
  | entropy_loss   | -0.809 |
  | explained_variance | 0.635 |
  | learning_rate  | 0.0001 |
  | loss           | 0.017 |
  | n_updates      | 390   |
  | policy_gradient_loss | 0.000601 |
  | value_loss     | 0.0452 |
  |-----|-----|
<stable_baselines3.ppo.ppo.PPO at 0x7f45ef77aa10>
```

## Images:

No signs of crashing in all 4 episodes



Graphs:

